

Introduction to Postgres and PostGIS

Databases: what are they?

- Usually when you hear the word “database” a person is referring to “Relational Databases” or “Relational Database Management Systems” (RDBMS)
- For the purposes of this talk, a database is: **“A system that allows for the creation and manipulation of tables of data”**
- Data in our tables can be references to rows in other tables (the “Relations” in our “Relational Database”)

Database Example: Hospitals

hospital

| id | name | telephone | general_beds | special_care_beds | address |
|----|--------------------|--------------|--------------|-------------------|--|
| 1 | SunnyTown Hospital | 123-555-1234 | 500 | 150 | 1234 Main Street, Townsville, Washington |
| 2 | Lady of the Lake | 123-555-2345 | 600 | 200 | 555 Street Rd, Cityville, Washington |
| 3 | Charity Hospital | 504-555-1234 | 2600 | 1000 | 1532 Tulane Ave, New Orleans, Louisiana |
| 4 | UW Medical Center | 206-123-4567 | 450 | 100 | 1959 NE. Pacific Street, Seattle, Washington |

department

| hospital_id | department | employees |
|-------------|-------------------------|-----------|
| 2 | Gastroenterology | 10 |
| 2 | Neurology | 5 |
| 2 | Primary Care | 50 |
| 3 | Primary Care | 40 |
| 3 | Rheumatology | 5 |
| 4 | Primary Care | 45 |
| 4 | Rehabilitation Medicine | 15 |

Structured Query Language: SQL

Examples

```
SELECT *  
FROM hospitals  
WHERE beds > 1000
```

| | | | | | |
|---|------------------|--------------|------|------|--|
| 3 | Charity Hospital | 504-555-1234 | 2600 | 1000 | 1532 Tulane Ave, New Orleans, Louisiana |
|---|------------------|--------------|------|------|--|

```
SELECT d.department  
FROM hospitals h  
INNER JOIN departments d ON h.id = d.hospital_id  
WHERE h.name = 'Lady of the Lake'
```

| |
|------------------|
| Gastroenterology |
| Neurology |
| Primary Care |

So where does GIS come in?

- Our simple database allows us to answer questions about Hospital's size and capabilities. But what about geospatial questions? Such as:
 - How many hospitals do I have in my County?
 - How many residents of this area are within 25 miles of a hospital?
 - Disaster strikes: How many hospitals do I have within 100 miles?


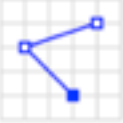
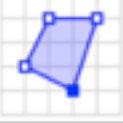

GIS Options

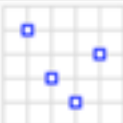

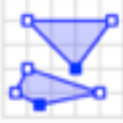
- We could parse out the City/State from our address, but this only allows us to count the number of hospitals per city/state.
- We could assign a latitude/longitude to each hospital. Issues with this:
 - We might have to write custom software to put on a map, off the shelf mapping options might not work.
 - Still need custom code for geospatial operations (Distance, Contains/Within/Intersects, etc...)
 - Could maybe have a solution for our Point data, but what about polygons and lines?

Enter PostGIS

- PostGIS adds support for geographic objects in Postgres: Points, LineStrings, Polygons and their “Multi” variations. Also supports Raster data (Digital Elevation Maps)
- Provides functions for determining geospatial interactions between objects (Within, Contains, Intersects) as well as measurements (Area, Length, Distance)
- Provides Spatial Indexes to vastly speed up geospatial operations
- Natively supported by off the shelf GIS tools: QGIS, Geoserver.

I stole this slide from wikipedia

| Type | Examples | |
|------------|---|--|
| Point |  | <code>POINT (30 10)</code> |
| LineString |  | <code>LINESTRING (30 10, 10 30, 40 40)</code> |
| Polygon |  | <code>POLYGON ((30 10, 40 40, 20 40, 10 20, 30 10))</code> |
| |  | <code>POLYGON ((35 10, 45 45, 15 40, 10 20, 35 10), (20 30, 35 35, 30 20, 20 30))</code> |

| Type | Examples | |
|-----------------|---|---|
| MultiPoint |  | <code>MULTIPOINT ((10 40), (40 30), (20 20), (30 10))</code> |
| | | <code>MULTIPOINT (10 40, 40 30, 20 20, 30 10)</code> |
| MultiLineString |  | <code>MULTILINESTRING ((10 10, 20 20, 10 40), (40 40, 30 30, 40 20, 30 10))</code> |
| MultiPolygon |  | <code>MULTIPOLYGON (((30 20, 45 40, 10 40, 30 20)), ((15 5, 40 10, 10 20, 5 10, 15 5)))</code> |
| | | <code>MULTIPOLYGON (((40 40, 20 45, 45 30, 40 40)), ((20 35, 10 30, 10 10, 30 5, 45 20, 20 35), (30 20, 20 15, 20 25, 30 20)))</code> |

Example geom column

| | name character varying(100) | admin character varying(200) | geom geometry(MultiPolygon,4326) |
|----|--------------------------------|---------------------------------|---|
| 1 | Aruba | Aruba | 0106000020E6100000010000000103000000010000001A000000362D7CD3CD7F51C0 |
| 2 | Badghis | Afghanistan | 0106000020E610000001000000010300000001000000030100005CAC4D5C99135040 |
| 3 | Hirat | Afghanistan | 0106000020E61000000100000001030000000100000060020000B006EA5F95AE4E40 |
| 4 | Khost | Afghanistan | 0106000020E610000001000000010300000001000000690000003487EB5832845140 |
| 5 | Pembroke | Bermuda | 0106000020E6100000010000000103000000010000000A000000EE397CD3ED3150C0 |
| 6 | Bamyan | Afghanistan | 0106000020E610000001000000010300000001000000810100009000CF2A9CEF5040 |
| 7 | Kapisa | Afghanistan | 0106000020E6100000010000000103000000010000004B000000AC52CFAA825B5140 |
| 8 | Balkh | Afghanistan | 0106000020E6100000010000000103000000010000006E010000502BF19295D05040 |
| 9 | Ordino | Andorra | 0106000020E61000000100000001030000000100000017000000800F04E9BD8EF93F0 |
| 10 | Canillo | Andorra | 0106000020E6100000010000000103000000010000001F00000080B3AF5A822DFC3F0 |
| 11 | Faryab | Afghanistan | 0106000020E610000001000000010300000001000000BB01000080E584AC0E5D5040 |
| | | | ----- |

Wall of text 1: Relationships and Measurements

- ST_Distance - Returns the distance between two objects (either 2D cartesian based on projection, or geodesic in meters depending on arguments)
- ST_Centroid - Returns geometric center of geometry
- ST_Disjoint - Returns TRUE if geoms don't intersect
- ST_Intersects - Returns TRUE if geoms intersect
- ST_Contains - Returns TRUE if one geom is inside another
- ST_Length - Returns length of geom (either 2D cartesian based on projection, or geodesic in meters depending on arguments)

Example 1: Distance Between Cities

```
SELECT ST_Distance(  
  (SELECT geom FROM geo.cities_ne WHERE name = 'Seattle'),  
  (SELECT geom FROM geo.cities_ne WHERE name = 'New Orleans'),  
  true) * 0.000621371 AS distance_miles
```

| | distance_miles double precision |
|----------|---|
| 1 | 2099.56714086922 |

Wall of text 2: Geometry Accessors/Editors

- ST_Envelope - Returns the bounding box of the geometry
- ST_SRID - Returns the Spatial Reference Identifier for the geometry
- ST_Transform - Returns a new Geometry with its coordinates transformed to the desired SRID
- ST_Affine - applies affine transformations to the geometry
- ST_Rotate - Rotates a geometry about an origin.

Example 2: Coordinate Information and Transformation

- First we will get the SRID of one of our geometries:

```
SELECT ST_SRID((SELECT geom FROM geo.cities_ne WHERE name = 'Seattle'))
```

| st_srid integer |
|--------------------|
| 4326 |

- Then we will transform it and observe the results

```
SELECT ST_AsText(  
  ST_Transform(  
    (SELECT geom FROM geo.cities_ne WHERE name = 'Seattle'),  
    3857  
  )  
)
```

| st_astext text |
|--|
| POINT(-13619041.444428 6035935.37594134) |

Wall of text 3: Geometry Constructors and Outputs

- ST_GeomFromText - Creates a geom from WKT
- ST_GeomFromKML - Takes KML input and creates a geom
- ST_AsGeoJSON - Returns geom as GeoJSON
- ST_AsText - Returns Well Known Text Representation of geom

Getting data into PostGIS

- One of the easiest ways to get data into PostGIS is via Shapefiles and the utility `shp2pgsql`
- If you have raster data you want to put into PostGIS (Digital Elevation Maps is the common use case) you can use `raster2pgsql`

Example 3: Furthest Cities

Which two cities on earth are furthest from each other?



Example 3: Furthest Cities

```
SELECT
    c1.name || ', ' || c1.adm0name AS city1,
    c2.name || ', ' || c2.adm0name AS city2,
    ST_Distance(c1.geom, c2.geom, true) * 0.000621371 AS distance_miles

FROM geo.cities_ne c1
CROSS JOIN geo.cities_ne c2
WHERE c1.scalerank < 3 AND c2.scalerank < 3 --Limit the run time of
                                           --the query by focusing on
                                           --bigger cities

ORDER BY ST_Distance(c1.geom, c2.geom, true) DESC
LIMIT 10
```

t pane

Output

Explain

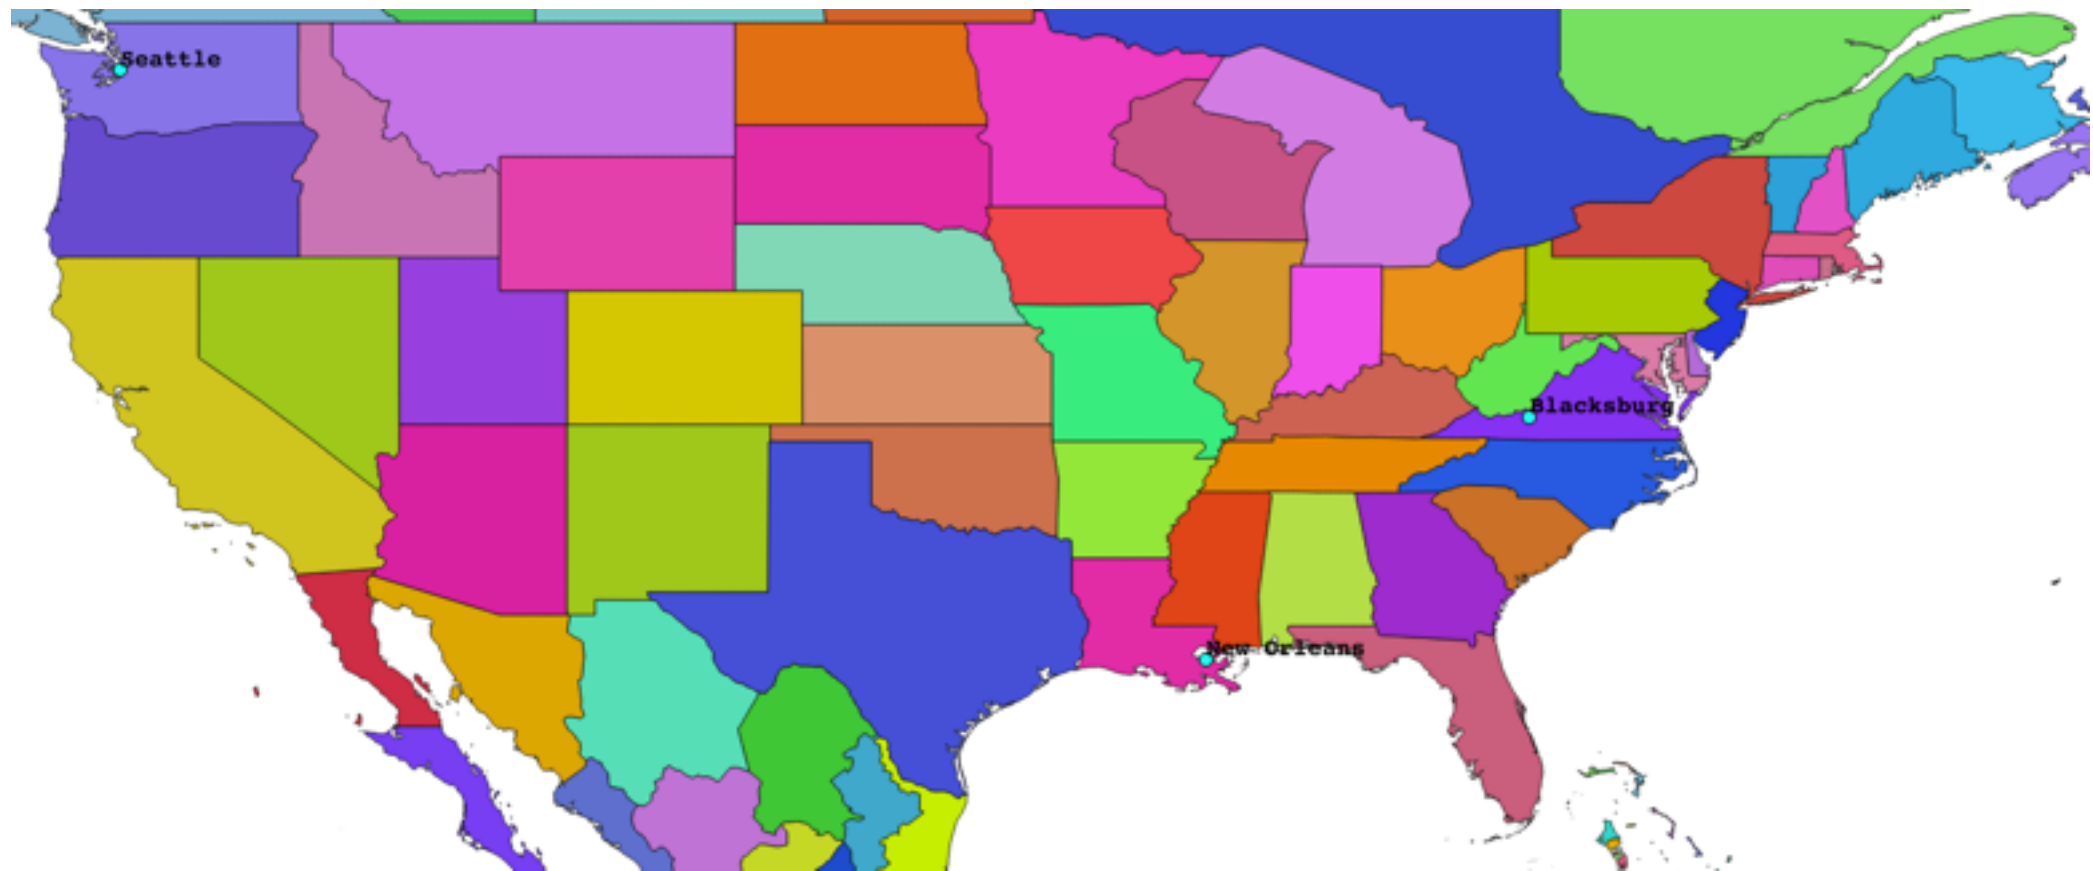
Messages

History

| city1 text | city2 text | distance_miles double precision |
|------------------------|------------------------|------------------------------------|
| Xian, China | Santiago, Chile | 12369.0959635262 |
| Santiago, Chile | Xian, China | 12369.0959635262 |
| Guayaquil, Ecuador | Kuala Lumpur, Malaysia | 12314.341883938 |
| Kuala Lumpur, Malaysia | Guayaquil, Ecuador | 12314.341883938 |

Example 4: Family Vacations

- My family lives all over the country, I'm in Seattle, my Sister is in Blacksburg Virginia, and my Parents are in New Orleans... Where is the fairest place for us to meet for Christmas?



Example 4: Solution

```
WITH fam_city AS(
    SELECT name, geom FROM geo.cities_ne c WHERE c.name = 'New Orleans'
    OR c.name = 'Blacksburg' OR c.name = 'Seattle'
)

SELECT c.name || ', ' || c.adm1name AS loc, ST_DISTANCE(c.geom,
    (SELECT ST_CENTROID(
        (SELECT ST_UNION(f.geom) FROM fam_city f))
    ), true) * 0.000621371 AS distance_miles

FROM geo.cities_ne c
ORDER BY ST_DISTANCE(c.geom, (SELECT ST_CENTROID(
    (SELECT ST_UNION(f.geom) FROM fam_city f)
)), true)
```

t pane

a Output

Explain

Messages

History

| loc text | distance_miles double precision |
|--------------------|------------------------------------|
| Hutchinson, Kansas | 22.4276893479084 |
| Salina, Kansas | 38.5301663778016 |
| Wichita, Kansas | 40.4335715748189 |

LAB TIME

- Go to: <http://maptimesea.github.io/2016/01/02/postgis-tutorial.html> To start the lab!
- Slides and QGIS file for viewing lab data can be downloaded at https://github.com/parkercoleman/postgis_talk (Click Download Zip)
- I can be reached at ADINSX.sa@gmail.com for questions :)